


Systematic Analysis of Second Order Optimization in Large-scale Neural Network Training

Linjian Ma*, Gabe Montague*, Jiayu Ye*,

Zhewei Yao, Amir Gholami, Kurt Keutzer, Michael W. Mahoney

University of California, Berkeley




Berkeley
UNIVERSITY OF CALIFORNIA



High-Level Outline

- Deep Learning and Large batch training
- Kronecker-Factored Approximate Curvature (K-FAC)
- Results: Level-playing-field performance comparison
- Results: Hyperparameter robustness behavior

Deep learning is the Trend



AlphaGo



Large Scale Deep Learning



Consuming data is slow

OPENAI FIVE

128,000 preemptible CPU cores on GCP

256 P100 GPUs on GCP

~180 years per day (~900 years per day
counting each hero separately)

Total: 10000 years

~2 realtime month

Large batch training saves time

- need fast training -> parallelization -> large batch


Objective function

$$L(\theta) = \frac{1}{N} \sum_{i=1}^N l(x_i, y_i, \theta)$$


Update rule

$$\theta_{t+1} = \theta_t - \eta_t \frac{1}{|B|} \sum_{(x,y) \in B} \nabla_{\theta} l(x, y, \theta_t)$$

small $|B|$





large $|B|$



Large batch SGD encounters problems

- Large batch SGD may get “stuck”:
 - Sharp minima
 - Saddle points
- Testing **and** training performance affected

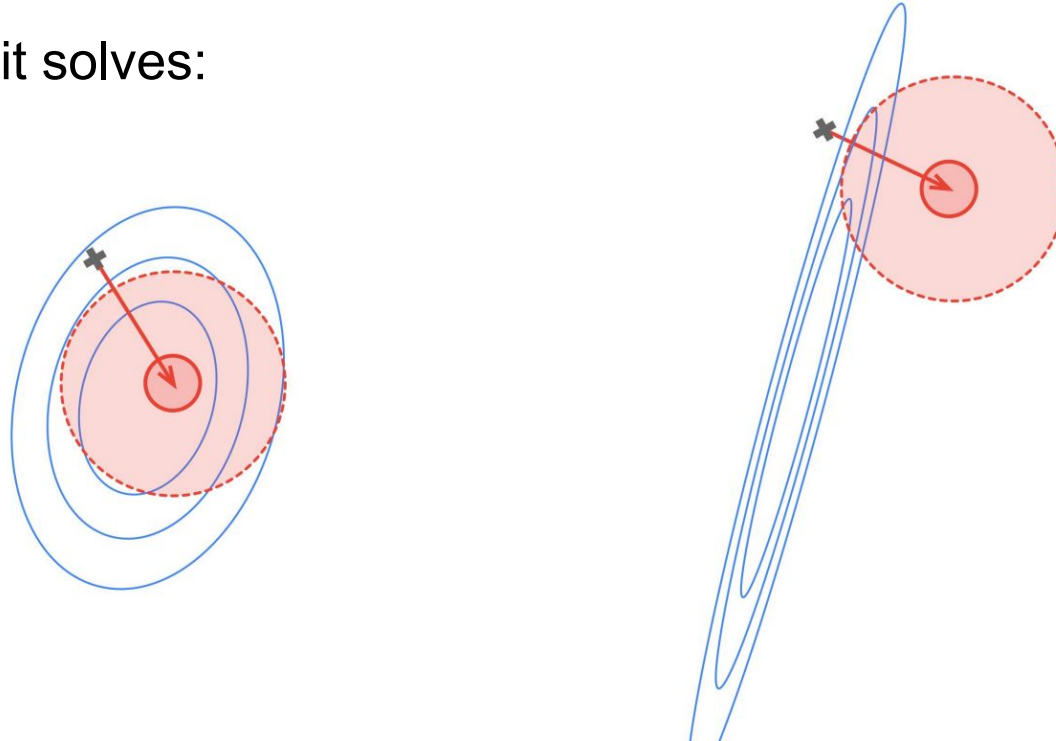


(a) Training without LARS

Loss landscape from <https://www.cs.umd.edu/~tomg/projects/landscapes/>
Ginsburg, Boris, Igor Gitman, and Yang You. "Large Batch Training of Convolutional Networks with Layer-wise Adaptive Rate Scaling." arxiv:1708.03888.

Second order method: K-FAC

- Kronecker-Factored Approximate Curvature (K-FAC)
- Problem it solves:

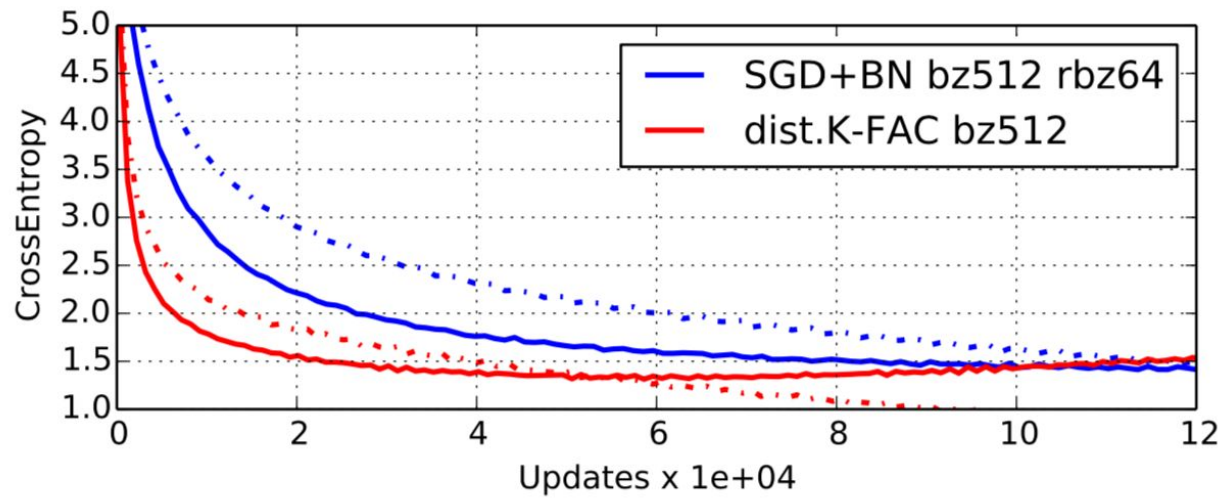


- Uses heavily-approximated second-order information to compute parameter updates

<https://supercomputersfordl2017.github.io/Presentations/K-FAC.pdf>

James Martens and Roger B. Grosse. "Optimizing Neural Networks with Kronecker-factored Approximate Curvature". In: *CoRR* abs/1503.05671 (2015). arXiv: 1503.05671.

Empirical studies of K-FAC vs. SGD



- Comparable performance to SGD
- Better for large batch training?

Training budget selection

- Different assumptions exist in literature:


Limiting factor is Time	Limiting factor is Money/Computation
<ul style="list-style-type: none">◦ Example: Owners of computing resources, e.g. Google◦ Easier hyperparameter tuning parallelization◦ Stop training based on number of updates: “iterations”	<ul style="list-style-type: none">◦ Example: Renters of computing resources, e.g. AWS users◦ Paying for each hyperparameter tried◦ Stop training based on number of epochs or training examples

A level-playing-field performance comparison

Hyperparameters of SGD	Hyperparameters of K-FAC
<ul style="list-style-type: none">◦ Learning rate◦ Momentum◦ Weight decay	<ul style="list-style-type: none">◦ Learning rate◦ Damping◦ Momentum◦ Weight decay◦ Clipping magnitude◦ Second-order update frequency◦ Second-order update momentum

- No relationships assumed between hyperparameters
- We tune two hyperparameters equally - giving the same amount of training to both methods



A level-playing-field performance comparison



- We race SGD and K-FAC to 3 target **accuracies**
- **Dotted: Ideal hyperbola** (linear because log-log)


A level-playing-field performance comparison

- Consistently slower = consistently worse end-of training performance (regardless of stopping rule)
- Our stopping rule:
 - Total training epochs = $(\log_2(\text{batch size}/128) + 1) \times 100$
- End-of-training (SGD vs. K-FAC):



Discussions

Did we get our results because we didn't tune carefully enough for K-FAC?



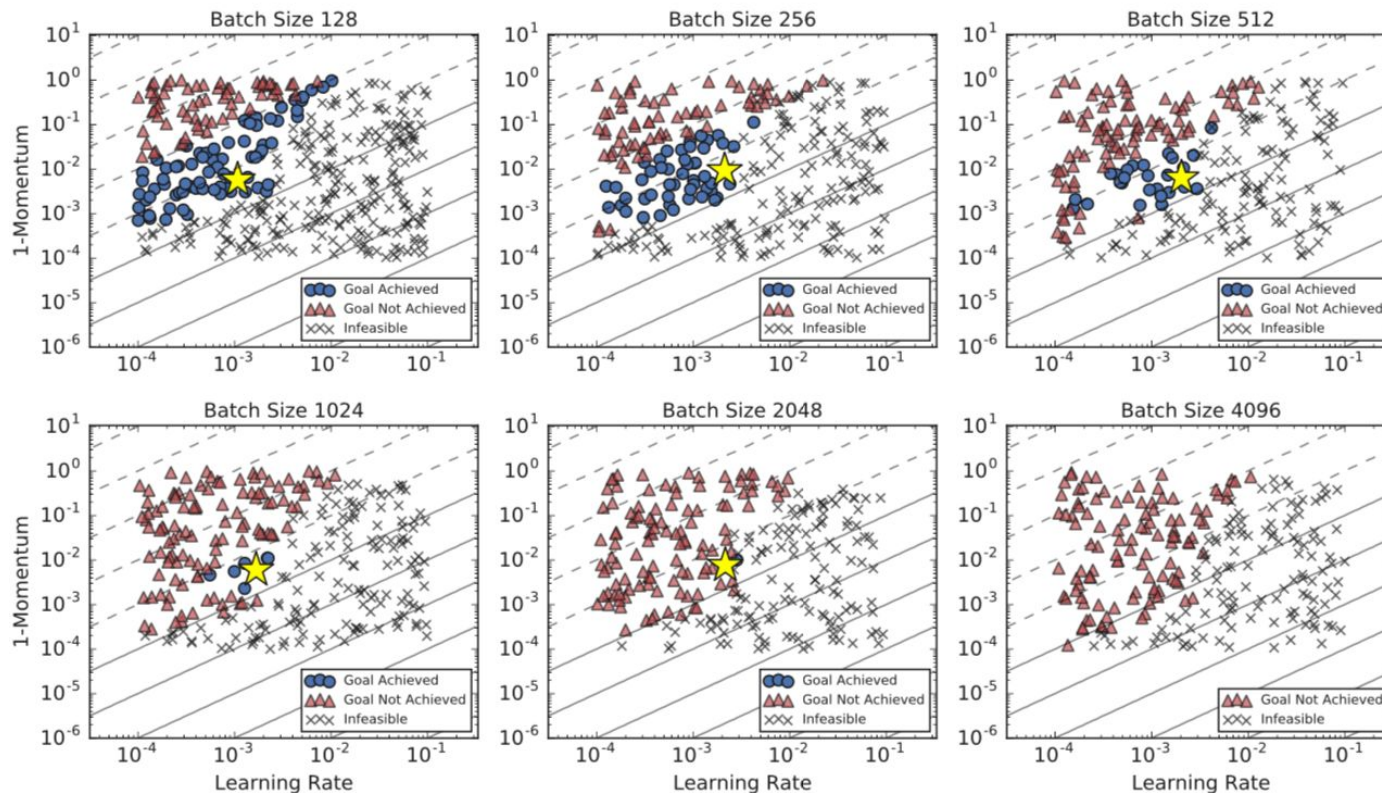
- **Hyperparameter tuning time is often ignored**
 - Justification: they are **“tune once” values** that **generalize**
- We used the “tune once” values for K-FAC and **it appears they did not generalize**

Conclusion: K-FAC hyperparameter tuning time = training time. We gave it the same amount as SGD

How do we measure hyperparameter robustness?

When limiting factor is money/computation (epochs):

Google SGD study:




Failure
Success

“Transformer on LM1B with a training budget of one epoch”

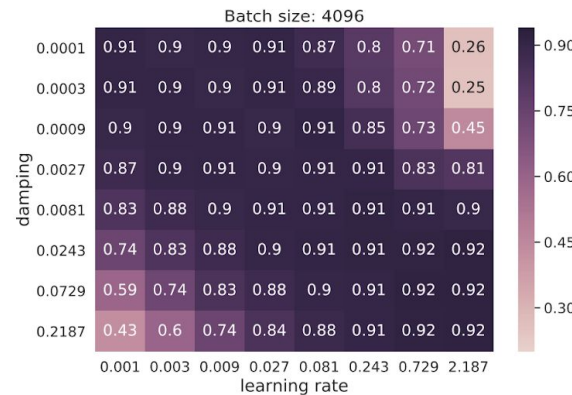
K-FAC Hyperparameter Robustness: Heatmaps

K-FAC Accuracy Heatmaps under our epoch-like budget:*


$|B| = 1,024$



$|B| = 4,096$



$|B| = 16,384$

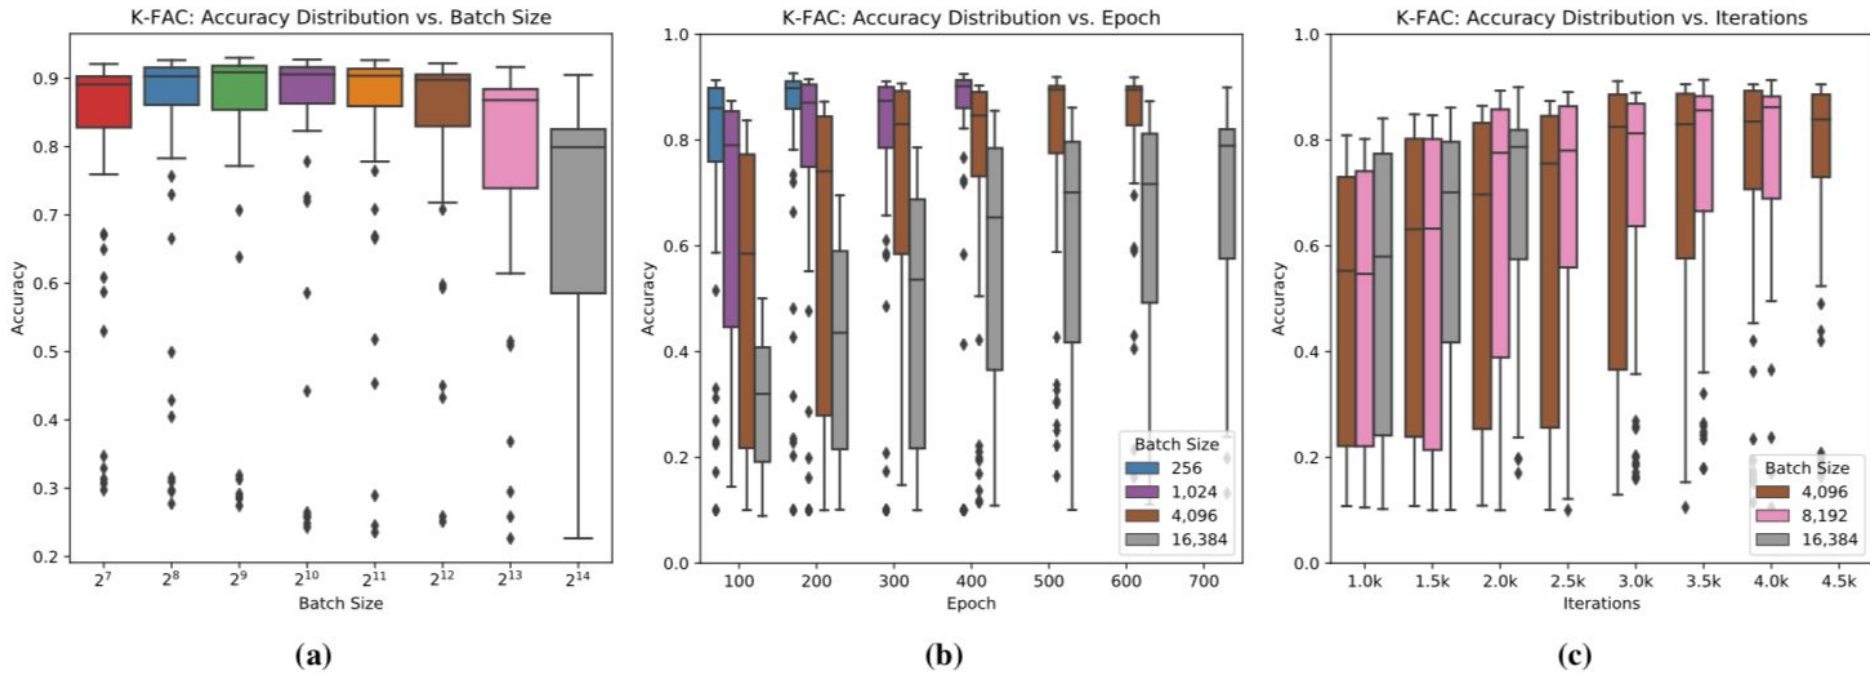


- **Shrinking** of the high-accuracy region with increasing batch size
- (Coincidental: positive correlation observed between damping and learning rate)

*Total training epochs = $(\log_2(\text{batch size}/128) + 1) \times 100$

K-FAC Hyperparameter Robustness: Box plots

K-FAC Accuracy Box Plots over time:



- **Growing** of the high-accuracy region with increasing batch size **with iteration budget** (not epoch budget)
- Aligns with Google finding for SGD
- Demonstrates competency of K-FAC algorithm

Summary of Empirical Results

- At small batch sizes, even with extensive hyperparameter tuning, K-FAC has **comparable**, but not superior, train/test performance to SGD.
- Increasing batch size for K-FAC results in **slower training** compared to SGD, when measured in terms of iterations.
- For fixed epochs, larger batch sizes result in **weaker** hyperparameter robustness.
- For fixed iterations, larger batch sizes result in **greater** hyperparameter robustness.

Summary of Findings

- Methods investigated: **K-FAC, TRCG (Trust Region Conjugate Gradient), Stochastic TRCG, SGD**
- *Similarities:*
 - General robustness patterns
 - Large-batch scalability problems
- *Differences:*
 - Large-batch performance, hyperparameter robustness